

OBGYN AI Scheduling Assistant

User Manual — Requirements & How-To Guide

FOR DEMONSTRATION ONLY — NOT FOR CLINICAL USE — SYNTHETIC DATA. All names, phone numbers (555-01xx) and MRNs (SYN-xxxx) are fictional. No PHI is present. The assistant is decision-support: it recommends, a human decides — nothing auto-executes.

This manual is organized around the client's **ASK list**. Each ASK is written as a **requirement** (REQ-1 ... REQ-6): the objective, the steps to achieve it (in the app and by chat), the expected result, and what runs behind the scenes. A final section covers the human-in-the-loop approval workflow that every change passes through.

Getting started

- 1 Open the assistant: <https://nychhc-frontend-iis-ai-ui.apps.ocp419.crucible.iisl.com>
- 2 Pick a role in the top-right role switcher. The role sets which tabs and actions you see: **Scheduler** (Selamawit — scheduling lead), **Approver** (Marcus — HR/operations), **Provider** (Dr. Chen), **Leadership** (Dr. Adeyinka — Chair/CCO).
- 3 Use the tabs to reach the work: **Dashboard** (KPIs + proactive insights), **Planning** (coverage, capacity, cancellations, template), **Approvals** (the HITL queue + audit trail), **Reporting** (department view, Leadership), and **Copilot** (chat).
- 4 Ask anything in **Copilot** in plain language — the assistant answers from the live data and remembers the conversation as you move between tabs.

Tip: every requirement below can be driven two ways — through the on-screen cards, or by typing the example prompts into the Copilot tab. Both use the same data and the same approval gate.

Requirements at a glance

Req	Ask	What it lets you do
REQ-1	ASK 1	Right-size the template — double-block & walk-in decisions from cancellation type
REQ-2	ASK 2	Approve PTO ahead without creating uncovered service windows
REQ-3	ASK 3	See the department as a whole and find the cycle-time bottleneck
REQ-4	ASK 4	Match providers to minute-weighted demand (headcount ≠ capacity)
REQ-5	ASK 5	Keep PHI in Epic — analyze aggregates, route patient-level work to Epic
REQ-6	ASK 6	Build the department-level value case for leadership

REQ-1

Appointment template & cancellation strategy (ASK 1)

Role: Scheduler

Objective

Decide whether to **double-block** a slot and whether to add a **walk-in provider** — driven by *why* appointments fall off. Advance cancellations free a slot you can refill; true no-shows waste it. The two need opposite responses.

How to — in the app

- 1 Go to the **Planning** tab.
- 2 Read the **Cancellations — advance vs true no-show** card. Each day/shift splits the cancel rate into *advance* (refillable) and *true no-show* (wasted).
- 3 Read the **Template recommendation** card: it flags “Double-block” only where true no-shows are high, and “Do NOT double-block” where cancels are mostly advance.
- 4 Check the walk-in line under Template for the Friday full-day vs half-day scenario and the projected hours/dollars saved.

How to — by chat (Copilot tab)

“How do cancellations break down by day this quarter?”

“Should we double-block Tuesday afternoons?”

“Do we need a full-day walk-in provider on Fridays?”

Expected result

Tuesday PM reads **advance-heavy** → *don't* double-block (tighten the waitlist instead). Monday AM reads **true-no-show-heavy** → *do* double-block. The Friday walk-in demand is AM-concentrated, so a **half-day** template covers it and frees ~48 provider-hours / ~\$5,280 a quarter with zero patients turned away.

Under the hood: Computed from the appointment-history corpus (each appointment tagged attended / advance-cancel / no-show). Deterministic and grounded — no guesswork.

REQ-2

Forward PTO coverage (approve-ahead) (ASK 2)

Role: Scheduler / Approver

Objective

Approve time off *before* it creates a problem. Coverage is a **skill-mix** question, not a headcount one — the assistant scans the next 90 days and tells you whether a leave request drops a service line below its minimum, and how to keep it covered.

How to — in the app

- 1 Open the **Planning** tab and read the **90-day coverage plan** card — days below a service-line minimum are listed with the providers out and the earliest at-risk date.
- 2 As **Approver**, open the **PTO** tab to see pending requests and their coverage impact.
- 3 Decide: approve, stagger the leave, or pull a peer/float onto service. Approving a request that would leave a window uncovered is **blocked** until explicitly overridden.

How to — by chat (Copilot tab)

“Where can't I cover service in the next 90 days?”

"Can I approve PTO for Dr. Wu July 14 to 18?"

Expected result

The plan ranks the gaps (e.g. the High-Risk Panel — Brooks + Wu — is the tightest). For a specific request the assistant answers approvable / not, and if not, offers concrete options (stagger the dates, add a per-diem). Final sign-off is always yours.

Under the hood: Coverage is evaluated against configurable service-line minimums over the roster + PTO calendar; the approval path is gated so nothing is silently approved into an uncovered window.

REQ-3

Consolidated department view (cycle time) (ASK 3)

Role: Leadership

Objective

See the department **as a whole** and find where time is actually lost. Cycle time (referral → seen) is the sum of hand-offs; the assistant attributes any slip to the specific stage rather than blaming 'capacity.'

How to — in the app

- 1 Switch to the **Leadership** role.
- 2 Open the **Reporting** tab and read the **cycle-time by hand-off** view — total days this period vs the prior quarter, broken into clerical intake, clinical scheduling, and provider availability.

How to — by chat (Copilot tab)

"How is the department performing as a whole?"

"Where's the cycle-time increase coming from?"

Expected result

Cycle time is up (~6.1 days vs ~5.4 last quarter) and the slip is isolated to the **clerical intake** hand-off — not provider capacity. The fix (consolidate intake / add logging capacity) compresses cycle time more than anything on the provider side.

Under the hood: Aggregated from per-referral hand-off timings (recent cohort vs prior); the bottleneck stage is identified by the largest period-over-period increase.

REQ-4

Capacity vs headcount (load balancing) (ASK 4)

Role: Scheduler / Approver

Objective

Match providers to **demand**, not to a head-count. A day can look balanced by bodies yet be over-loaded once you weight by visit-type minutes (a New-OB visit is far longer than a follow-up).

How to — in the app

- 1 Open the **Planning** tab and read the **Provider capacity — minute-weighted** card: per-weekday utilization with an over-loaded / under-utilised flag.
- 2 Read the one-line rebalance recommendation (e.g. "move one provider" from the slack day to the over-loaded day).
- 3 Click "**Draft proposal** →" (top-right of the card). The rebalance is staged into the **Approvals** queue and you're taken straight there.

How to — by chat (Copilot tab)

"Is our provider distribution actually matching demand?"

"Submit it to the approver as a proposal"

Expected result

Utilization differs by weekday even when head-count looks even; the over-loaded day is flagged with a cost-neutral rebalance. "Draft proposal" (or the chat command) places it in the approver queue for sign-off — it does not change the schedule by itself.

Under the hood: Demand per weekday comes from a **forecast model served by KServe on OpenShift AI** (nychhc-forecast), divided by staffed minutes; the no-show risk model is served the same way. If a model is unreachable the view falls back to history and says so.

REQ-5

Epic as system of record & PHI boundary (ASK 5)

Role: Scheduler / Approver

Objective

Keep **Epic the single source of truth** and keep PHI out of the assistant. It analyzes de-identified aggregates and sends any patient-level, actionable output back into Epic — so nothing lives in two places.

How to — in the app

- 1 Work from the aggregate cards (coverage, capacity, cancellations) as usual — these carry no patient identifiers.
- 2 When you need to act on specific patients, use the assistant to post an aggregate alert to the Epic scheduling chat; the named work happens in Epic.

How to — by chat (Copilot tab)

“Which specific patients are affected on Sept 9?”

“Post a coverage alert to the Epic chat”

“Where are you getting these numbers?”

Expected result

A request for a named patient list is **declined here and routed to Epic** (the assistant keeps it at the aggregate). Posting an alert writes the date/service/options to the Epic scheduling chat and is recorded in the audit trail. “Where are you getting this?” explains that all data is read from Epic.

Under the hood: The assistant reaches data only through an Epic/FHIR adapter seam — it never holds Epic credentials. Posting an alert is a role-gated, audited action (Scheduler/Approver).

REQ-6

Department-level value case (ASK 6)

Role: Scheduler / Leadership

Objective

Turn the day-to-day wins into a **department-level justification** a Chair/CCO will weigh — framed as a department outcome, not one director's preference.

How to — in the app

- 1 Gather the supporting numbers from Planning (capacity, cancellations) and Reporting (cycle time).
- 2 Ask the assistant to assemble the case; review and edit before sharing.

How to — by chat (Copilot tab)

“Help me make the case for the chair”

“Build the one-page business case for OBGYN scheduling”

Expected result

A short, department-framed summary built from the real metrics: the cost-neutral rebalance, the half-day Friday walk-in savings, the cycle-time fix at intake — each stated as a department win with the number behind it.

Under the hood: Synthesized from the same grounded analytics used elsewhere; advisory — you confirm before acting or sharing.

The approval workflow (every change passes through here)

Nothing the assistant proposes is applied automatically. Every recommendation becomes a **proposal** that a human approves, modifies, or rejects — and every decision is written to an audit trail with the user and timestamp.

Walk-through

- 1 Draft** — from the capacity card click “Draft proposal→”, or tell the assistant “submit it as a proposal.” It is staged, not applied.
- 2 Find it** — open the **Approvals** tab (Scheduler/Approver). The proposal appears in the pending queue with its summary and rationale.
- 3 Decide** — **Approve, Modify, or Reject**. A change that would leave a service window uncovered is blocked unless explicitly overridden.
- 4 Audit** — the decision is recorded (who, what, when, outcome) and visible in the audit trail on the same tab.

Why this matters: it is the guardrail behind every ASK — the assistant does the analysis and drafts the action; a named human always makes the final call, and the record proves it.

Good to know

- **Synthetic data only.** Names, phones (555-01xx) and MRNs (SYN-xxxx) are fictional; no PHI.
- **AI models run on OpenShift AI / KServe.** Two models are served and used: no-show risk and demand forecast. The conversational assistant uses Claude; the on-screen answers are grounded in the live data.
- **Session memory.** The Copilot conversation persists as you move between tabs and roles — it does not reset. Use “Clear” to start fresh.
- **Roles gate actions.** Providers can view their own schedule but cannot approve or post; Leadership sees the department reporting view.